



# SQL Workshop

A3SR Stats Club



# Getting Started

Download DBeaver v5.2.4

Download chinook.db

Load the Data in

# Intro

## Relational Databases:

A database structured to recognize relations among stored items of information.

## Other types:

- Non-relational Databases
- Graph Databases
- Etc.



# Why use databases?

## **Save space:**

- Optimized data structure (schema)
- Facts Tables:
  - events
- Dimensions Tables:
  - attributes

## **Save Time:**

- Indexing
  - Filter (subsetting)
  - Aggregates (grouping)
  - Joins (merging)

# Terminology

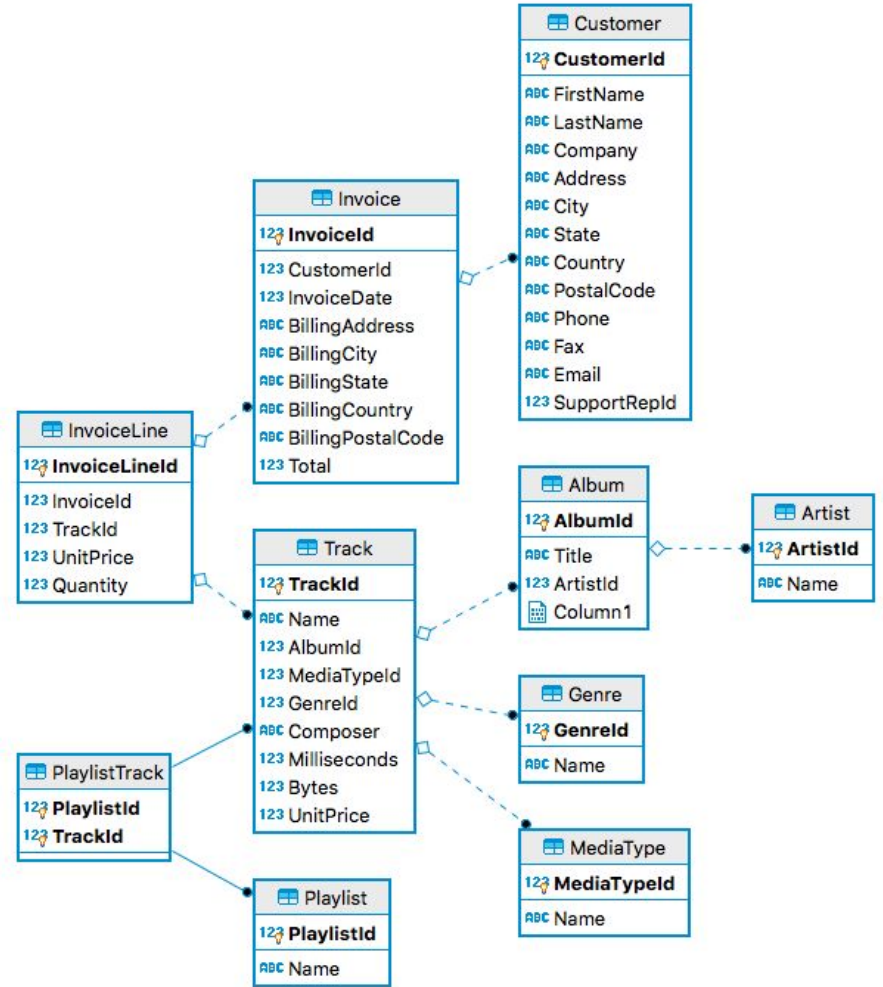
Schema

Tables

Columns

Primary Key (PK)

Foreign Key (FK)



# Terminology

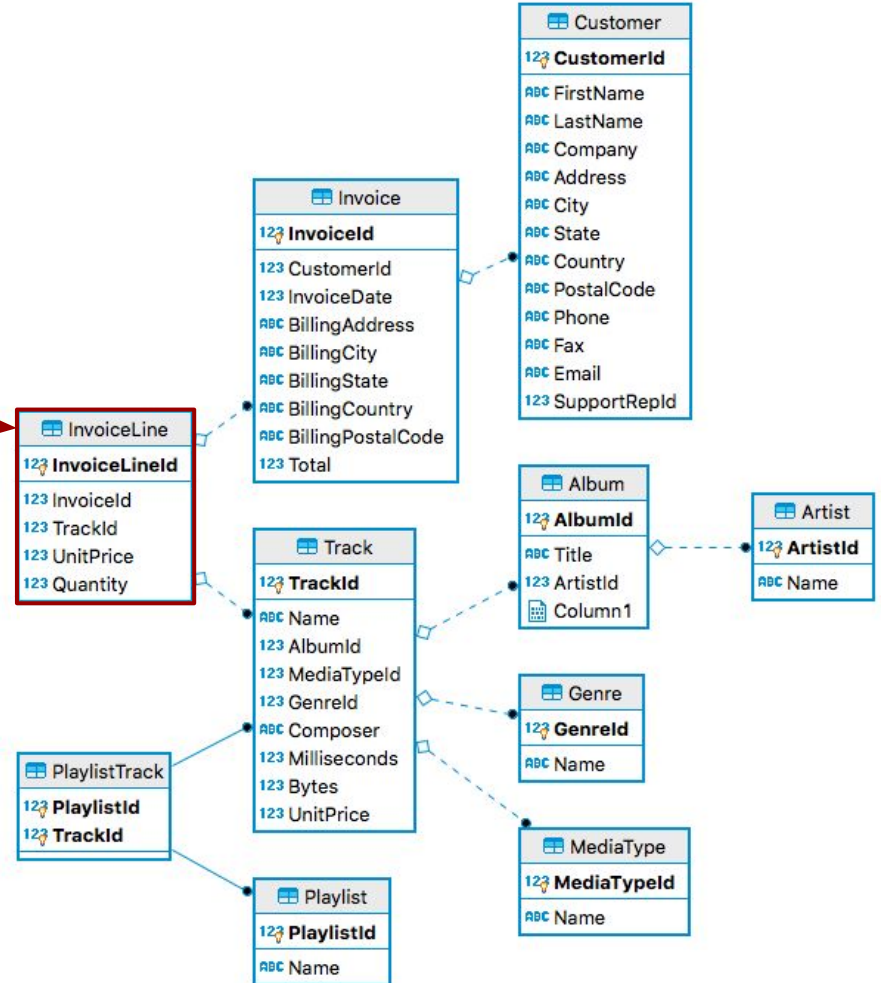
Schema

Tables

Columns

Primary Key (PK)

Foreign Key (FK)



# Terminology

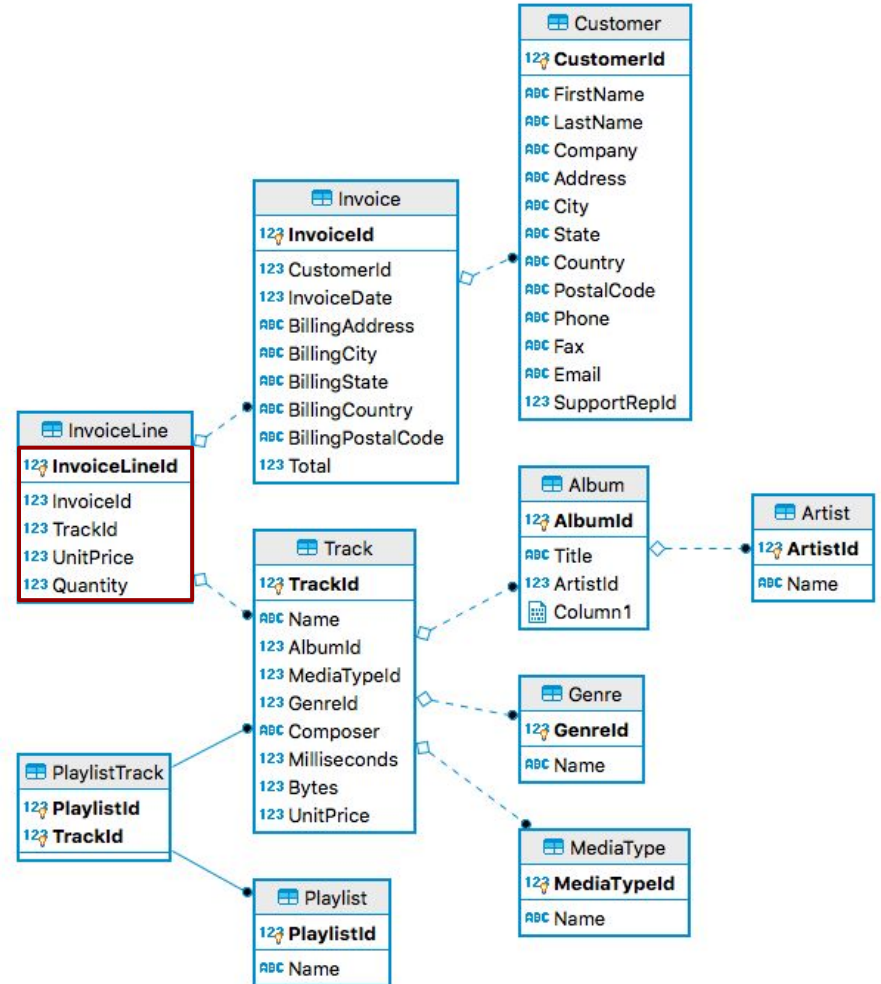
Schema

Tables

Columns

Primary Key (PK)

Foreign Key (FK)



# Terminology

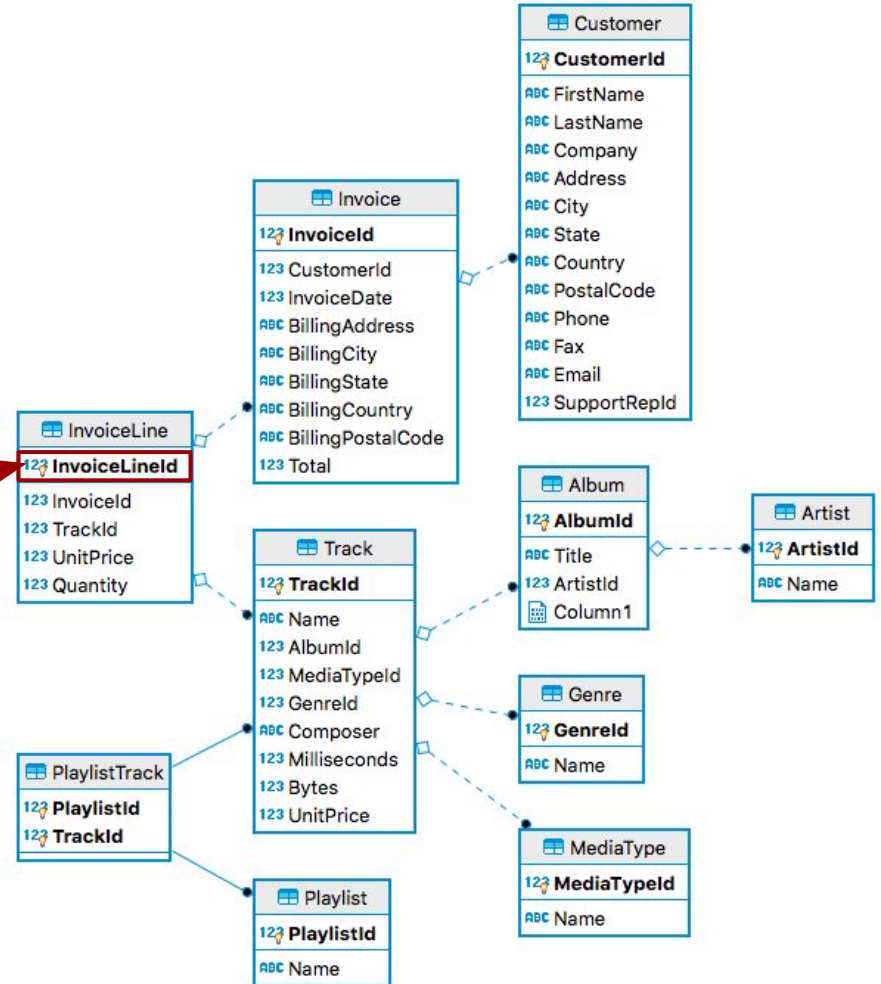
Schema

Tables

Columns

Primary Key (PK)

Foreign Key (FK)





# Terminology

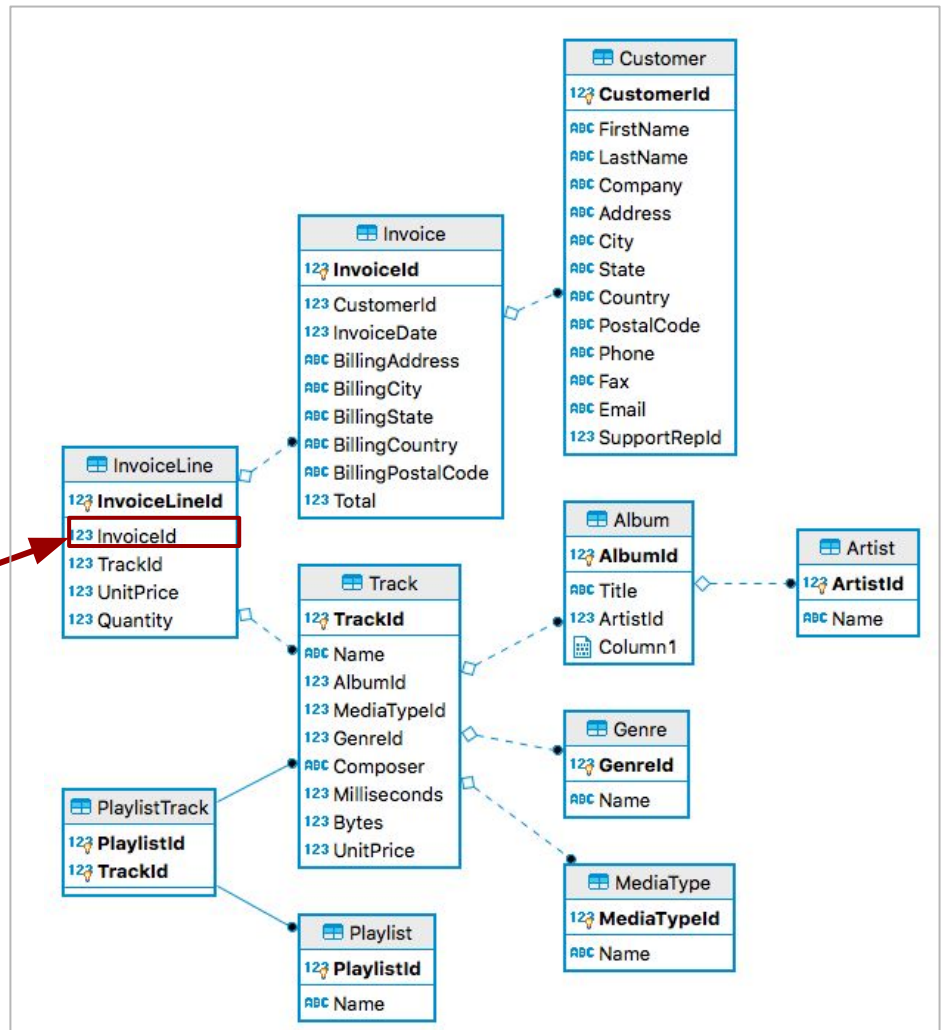
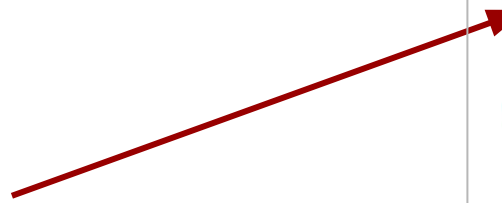
Schema

Tables

Columns

Primary Key (PK)

Foreign Key (FK)



# Syntax : 'Querying' the tables

SELECT *columns* FROM *table* ;

SELECT \* FROM *table* ;

SELECT \* FROM *table* *t* ; (*t*: Alias for the query)

SELECT \* FROM *table* LIMIT 5 ;

SELECT \* FROM *table* ORDER BY *column* ;

SELECT \* FROM *table* ORDER BY *column* DESC ;

# Syntax: 'Filtering' the table

SELECT *columns* FROM *table* WHERE *column* = x ;

SELECT *columns* FROM *table* WHERE *column1* = x AND *column2* < y ;

SELECT DISTINCT *column* FROM *table* ;

SELECT *columns* FROM *table* WHERE *column* BETWEEN x AND y ;

## Notes:

Conditional Operators: = , != , > , < , >= , <=, etc.

Casting: CAST(*column* as datatype) [eg: INT, CHAR]

# Syntax: 'Aggregating' the table

SELECT COUNT(\*) FROM *table* ;

Counts the rows

SELECT AVG(*column*) FROM *table* ;

Average value

SELECT SUM(*column*) FROM *table* ;

Total value of the rows

SELECT MIN(*column*) FROM *table* ;

Min value in the column

SELECT MAX(*column*) FROM *table* ;

Max value in the column

# Syntax: 'Grouping' by values in a column

```
SELECT column, COUNT(*) FROM table GROUP BY column ;
```

```
SELECT column, COUNT(*) FROM table GROUP BY column HAVING column > x ;
```

Exactly the same idea as dplyr!

# Order of Operations: Structuring the query

SELECT	<i>some columns, aggregates</i>
FROM	<i>some table</i>
WHERE	<i>condition(s) satisfied</i>
GROUP BY	<i>variables in a column</i>
HAVING	<i>some aggregate condition satisfied</i>
ORDER BY	<i>a variable</i>
LIMIT	<i>the results</i>

# Exercise-1a

1. How many tables are in your Schema?
2. What are the columns of the InvoiceLine Table?
3. Now looking at the Track table: what are the Primary Keys and Foreign?
4. What are the 3 shortest tracks (in milliseconds)?
5. Let's look at songs longer than 300000 ms. What is the title of the song with the smallest bytes?

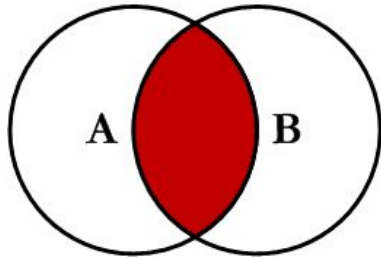
# Exercise-1b

6. Let's look at the Invoice table. What is the total number of orders?
7. How many orders were made in 2008? Use 'between'
8. Let's look at the Customer table. What is the number of unique countries that customers come from?
9. How many customers come from Canada?
10. From the Tracks table, find the GenreID corresponding to the Genre with the highest average track length accounting only those tracks greater than 30secs. With this information of the GenreID, find the name of the Genre from the Genre table.

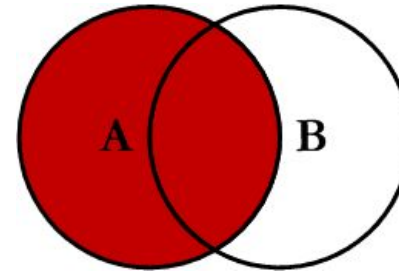


# Joins

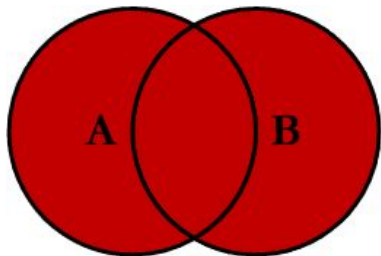
JOIN (INNER JOIN)



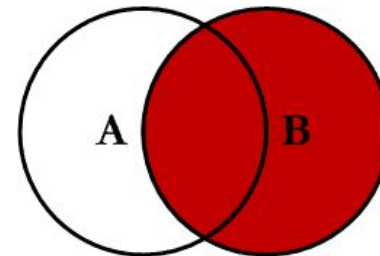
LEFT JOIN



OUTER JOIN



RIGHT JOIN



# Syntax: 'Joining' multiple tables

```
SELECT column_table1, column_table2 FROM table1 JOIN table2  
USING(primary_key);
```

```
SELECT column_table1, column_table2 FROM table1 JOIN table2 ON  
primary_key_table1 = foreign_key_table2
```

```
SELECT column_table1, column_table2 FROM table1 OUTER JOIN table2  
USING(primary_key);
```

```
SELECT column_table1, column_table2 FROM table1 RIGHT JOIN table2  
USING(primary_key);
```

# Exercise-2

1. Which album by Metallica is divided in two discs ?
2. How many songs are in each of the top 3 playlist (provide names of playlists)?
3. How much did Martha Silk spend on songs in total?
4. How many playlists include songs by Iron Maiden?

# Advanced SQL

Views: WITH ( ... ) AS *xyz*

Window Functions:

Aggregate to Arrays:

Geospatial querying with PostGIS & PostgreSQL:

# Best Practices

**Indexing:** Speeds up querying immensely on the columns that you expect to filter and group by often

**Order of Joins:** Smallest to largest. Use views to filter before joining.

**Test before you run:** Use limit to restrict size to test your query on before running it on the full table

# Order of Execution (behind the scenes)

ORDER	CLAUSE	FUNCTION
1	from	Choose and join tables to get base data.
2	where	Filters the base data.
3	group by	Aggregates the base data.
4	having	Filters the aggregated data.
5	select	Returns the final data.
6	order by	Sorts the final data.
7	limit	Limits the returned data to a row count.

# HAPPY BDAY GEORGE

